

Visualizing Deep Learning using PyTorch and PowerAI

Dustin VanStee / Bob Chesebrough
Data Scientists
IBM Worldwide Client Experience Centers

2019 IBM Systems Technical University
Oct 7-11 | Las Vegas, NV



IBM Systems Worldwide Client Experience Centers



IBM Systems Worldwide Client Experience Centers

maximize IBM Systems competitive advantage in the Cloud and Cognitive era by providing access to world class **technical experts** and **infrastructure services** to assist Clients with the transformation of their IT implementations..

9 Worldwide Locations (* also Infrastructure Hubs):

Austin TX , *Poughkeepsie NY, Rochester MN,
Tucson AZ, *Beijing CHINA, Boeblingen GERMANY,
Guadalajara MEXICO,*Montpellier FRANCE, Tokyo
JAPAN



Client Experience	Architecture & Design	Infrastructure Solutions	Content
Tailored, in-depth technology Innovation Exchange Events Relationship building Demonstrations Meetups Solution workshops Remote options (Inbound & Outbound)	Advise clients, "Art of the Possible" Discovery & Design Workshops, Consulting, Showcases, Reference Architectures, Co-Creation of assets (Inbound & Outbound)	Benchmarks, MVP & Proof of Technology "Test Drives" Demonstrations Infrastructure Services Certify ISV solutions Hosting Cloud Environment (Inbound to Centers)	Content Development IBM Redbooks Training Courses Video courses "Test Drives" Demonstrations

NEW: Co-Creation Lab; CEC Cloud; RedHat Center of Competency

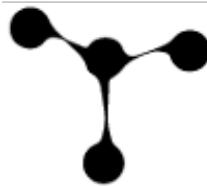
For further information, please contact the Centers via email at: ccenter@us.ibm.com

Agenda

- Pytorch on PowerAI
 - What is Pytorch?
 - Where does it fit in deep learning frameworks
 - Where can I get it?
 - Some examples
 - GPU access
 - Demo
 - Lab !

Deep Learning Frameworks

Caffe

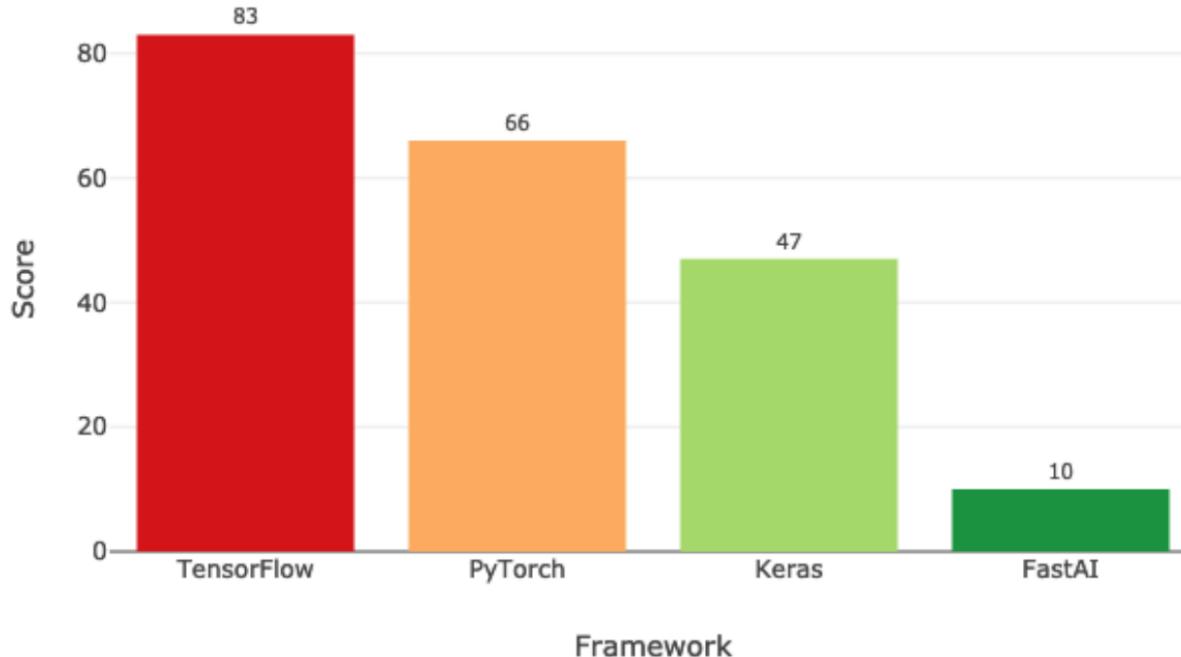


TensorFlow



Deep Learning Framework Popularity

Deep Learning Framework Six-Month Growth Scores 2019



- <https://towardsdatascience.com/which-deep-learning-framework-is-growing-fastest-3f77f14aa318>

PYTORCH – a quick review

Facebook's framework for research

- ✓ Cousin of LUA based Torch framework, but was rewritten to be tailored to Python frontend
- ✓ Gaining popularity quickly for its ease of use in R&D
- ✓ Supports dynamic computation graphs !
- ✓ Based on Python with Numpy compatibility
- ✓ Multi-GPU
- ✓ Easy to use, and supports standard debug tools



PyTorch Build	Stable (1.0)				Preview (Nightly)	
Your OS	Linux				Mac	
Package	Conda				LibTorch	
Language	Python 2.7	Python 3.5	Python 3.6	Python 3.7	C++	
CUDA	8.0	9.0	10.0	None		
Run this Command:	<code>conda install pytorch torchvision cudatoolkit=10.0 -c pytorch</code>					

PyTorch Basics – ND arrays

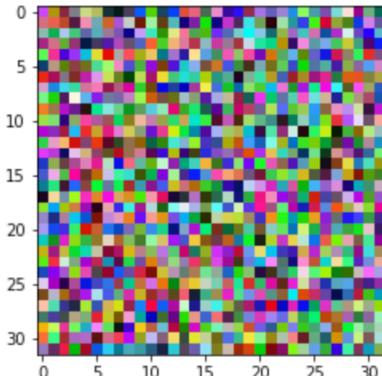
ND Tensors

When working with neural networks, you are always dealing with multidimensional arrays. Here are some quick tricks

Assume A is a 32x32 RGB image

```
## 3D Tensors
import torch
A = torch.rand(32,32,3)
plt.imshow(A)
```

<matplotlib.image.AxesImage at 0x10007f7ac208>



PyTorch Basics – Basic Slicing /Dimension Operations

Slicing Tensors - grab 'RED' dimension

```
In [15]: red_data = A[:, :, 0] #0 represents the first channel of RGB  
red_data.size()  
  
Out[15]: torch.Size([32, 32])
```

Swap the RGB dimension and make the tensor a 3x32x32 tensor

```
In [11]: A_rgb_first = A.permute(2, 0, 1)  
print(A_rgb_first.size())  
  
torch.Size([3, 32, 32])
```

Add a BatchSize to our Image Tensor

Usually you need to do this to run inference on your trained model

```
In [20]: Anew = A.unsqueeze(0)  
print(Anew.size())  
  
torch.Size([1, 32, 32, 3])  
torch.Size([32, 32, 3])
```

PyTorch Basics – Basic Matrix Operations

Matrix Multiply

```
In [18]: a = torch.randint(4,(2,3))
b = torch.randint(4,(3,2))
print(a)
print(b)
```

```
tensor([[0, 0, 1],
        [0, 2, 1]])
tensor([[1, 1],
        [3, 3],
        [0, 1]])
```

```
In [19]: # 2x3 @ 3x2 ~ 2x2
a.matmul(b)
torch.matmul(a,b)
```

```
Out[19]: tensor([[0, 1],
                  [6, 7]])
```

PyTorch Basics – Basic Index Operations

Create a onehot vector

In [40]:

```
batch_size = 5
nb_digits = 10
# Dummy input that HAS to be 2D for the scatter (you can use view(-1,1) if needed)
y = torch.LongTensor(batch_size,1).random_() % nb_digits
# One hot encoding buffer that you create out of the loop and just keep reusing
y_onehot = torch.FloatTensor(batch_size, nb_digits)

# In your for loop
y_onehot.zero_()
y_onehot.scatter_(1, y, 1)

print(y)
print(y_onehot)

tensor([[ 8,
         1,
         4,
         5,
         7]])
tensor([[ 0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.],
        [ 0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.],
        [ 0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  0.],
        [ 0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  0.],
        [ 0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,  0.]])
```

PyTorch: Using GPU's

CUDA Tensors

Tensors can be moved onto any device using the `.to` method.

```
In [17]: # let us run this cell only if CUDA is available
# We will use ``torch.device`` objects to move tensors in and out of GPU
x = torch.rand(2,2,2)
if torch.cuda.is_available():
    device = torch.device("cuda")           # a CUDA device object
    y = torch.ones_like(x, device=device)   # directly create a tensor on GPU
    x = x.to(device)                      # or just use strings ``.to("cuda")``
    z = x + y
    print(z)
    print(z.to("cpu", torch.double))       # ``.to`` can also change dtype together!
```

```
tensor([[ [1.5508, 1.0580],
          [1.6207, 1.1363]],

         [[1.9148, 1.6978],
          [1.5459, 1.5224]]], device='cuda:0')
tensor([[ [1.5508, 1.0580],
          [1.6207, 1.1363]],

         [[1.9148, 1.6978],
          [1.5459, 1.5224]]], dtype=torch.float64)
```

PyTorch: Build a model

To build a neural network using pytorch, you need to perform three main steps

- extend `nn.Module`
- define `init` function
- define forward function

```
In [ ]: class NetCNN3L(nn.Module):  
  
    def __init__(self):  
        super(NetCNN3L, self).__init__()  
        self.name = "NetCNN3L"  
        # output dimension (H,W) -> H - kernel + 1 - 2p  
        # in_NCHW=[Nx3x32x32 image], 3x3 square kernel, out_NCHW=[Nx32x32x32 image]  
        # cin=3,cout=32  
        self.conv1_1 = nn.Conv2d(3, 32, kernel_size=(3,3),padding=(1,1)) # same padding  
        # in_NCHW=[Nx3x32x32 image], 3x3 square kernel, out_NCHW=[Nx32x32x32 image]  
        self.conv1_2 = nn.Conv2d(32, 32, kernel_size=(3,3),padding=(1,1))# same padding  
        # in_NCHW=[Nx32x16x16 image], 3x3 square kernel, out_NCHW=[Nx64x16x16 image]  
        self.conv2_1 = nn.Conv2d(32, 64, kernel_size=(3,3),padding=(1,1))  
        # in_NCHW=[Nx64x16x16 image], 3x3 square kernel, out_NCHW=[Nx64x16x16 image]  
        self.conv2_2 = nn.Conv2d(64, 64, kernel_size=(3,3),padding=(1,1))  
  
        # an affine operation: y = Wx + b  
        # 64 x 8 x 8  
        self.fc1 = nn.Linear(4096, 512)  
        self.fc2 = nn.Linear(512, 10)  
  
    def forward(self, x):  
        # Max pooling over a (2, 2) window  
        x = self.conv1_1(x); self.c1_1 = x # (for plotting)  
        x = F.relu(x)  
        x = self.conv1_2(x); self.c1_2 = x # (for plotting)  
        x = F.relu(x)  
        x = F.max_pool2d(x, (2,2))  
        x = F.dropout(x, p=0.25)  
        x = self.conv2_1(x); self.c2_1 = x # (for plotting)  
        x = F.relu(x)  
        x = self.conv2_2(x); self.c2_2 = x # (for plotting)  
        x = F.relu(x)  
        x = F.max_pool2d(x, (2,2))  
        x = F.dropout(x, p=0.25)  
        #Flatten x for fully connected layer  
        x = x.view(-1, 4096)  
        #print(x.size())  
        x = self.fc1(x)  
        #print(x.size())  
        x = F.relu(x)  
        x = F.dropout(x, p=0.5)  
        x = self.fc2(x)  
        x = F.softmax(x)  
        return x  
  
    def num_flat_features(self, x):  
        size = x.size()[1:] # all dimensions except the batch dimension  
        num_features = 1  
        for s in size:  
            num_features *= s  
        return num_features
```

PyTorch: Train a model

```
for i,(X,Y) in enumerate(data_loader):
```

The key step during training is that after running an inference on a batch

```
yhat = model(X)
```

We calculate the loss

```
loss = criterion(yhat, Y)
```

and then update the weights based on the calculated gradients

```
loss.backward()
optimizer.step()
```

All the rest of this function is just mainly used for book keeping ...

PyTorch: Inference on a model

```
In [22]: # Send x to the GPU and run inference
# Select an image from our test set
IMG = 17
x = torch.tensor(np.asarray(test_set[IMG][0]), dtype=torch.float32)

# Plot the image (requires H,W,C as input so need to permute the axes)
plot_image(x.permute(1,2,0))

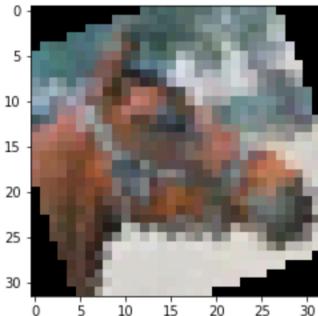
# For inference, convert C,H,W to N,C,H,W by adding the batch dimension . Batch size=1
x = x.unsqueeze(0) # add an batch dimension

y = model(x.to('cuda'))
y

torch.Size([32, 32, 3])

/gpfs/gpfs_g14_16mb/s4s004/vanstee/anaconda3/envs/powerai-1.6.0/lib/python3.6/site-packages/invkernel_launcher.py:59:
UserWarning: Implicit dimension choice for softmax has been deprecated.
```

```
Out[22]: tensor([[0., 0., 0., 0., 0., 0., 0., 1., 0., 0.]], device='cuda:0',
grad_fn=<SoftmaxBackward>)
```



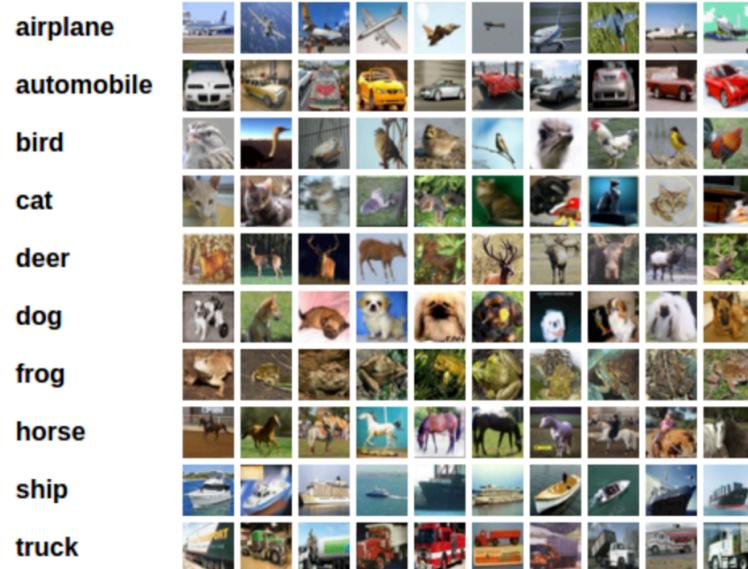
```
In [6]: labels_map = {
    0: "plane",
    1: "car",
    2: "bird",
    3: "cat",
    4: "deer",
    5: "dog",
    6: "frog",
    7: "horse",
    8: "ship",
    9: "truck"
}
```

PyTorch : TorchVision Library

Torchvision : An image library that contains

- Popular Image Datasets
- Popular NN Models (pretrained)
- Transformers

- **AlexNet**
- **VGG**
- **ResNet**
- **SqueezeNet**
- **DenseNet**
- **Inception v3**



CIFAR 10 Dataset

PyTorch : Vision Dataset API Example



```
xt = transforms.Compose([
    #transforms.CenterCrop(10),
    transforms.ToTensor(),
])
```

```
train_fm_set = dset.FashionMNIST(root = './fashion_mnist_data',
                                   download=True, transform=xt)
test_fm_set = dset.FashionMNIST(root='./fashion_mnist_data', train=False,
                                 download=True, transform=xt) #transform=trans
```

```
| test_fm_loader = torch.utils.data.DataLoader(
|     dataset=test_fm_set,
|     batch_size=batch_size,
|     shuffle=False)
```

Torchvision : An image library that contains

- Popular Image Datasets
- Popular NN Models (pretrained)
- Transformers



Pytorch PowerAI LAB –

Dustin VanStee

Bob Chesebrough

IBM WW Client Centers

Lab Code Here :

Code is here

<https://github.com/dustinvanstee/pytorch-examples>

